

Tracking-Data Driven Tactical Identity and Decision Support System in Professional Football (Soccer)

Project by Anthony Schauer

System Architecture / Workflow Specification

System Overview

This project is designed as a layered analytics system rather than a single analysis script or report. The system combines a core football tracking-data analytics pipeline with an AI-assisted workflow layer that operates on top of computed outputs. The purpose of this design is to preserve analytical rigor while also demonstrating how modern AI tools can be integrated into real analytical workflows to improve interpretation, communication, and decision support.

The architecture is intentionally built around the principle that **AI should not replace the analysis itself**. Instead, it should assist where human analysts often spend large amounts of time translating quantitative findings into structured qualitative outputs, recommendations, and reports. This makes the project both analytically credible and professionally relevant to emerging AI-integrated analytical roles.

The full system can be understood as five connected layers:

1. data ingestion and preparation,
2. feature engineering and tactical metric computation,
3. tactical modeling and identity analysis,
4. AI-assisted interpretation and decision-support generation,
5. reporting and output delivery.

System Design Philosophy

This architecture is based on four core design principles.

1. Analytical outputs come first

The foundation of the system is the tracking-data analysis itself. All downstream interpretation must depend on computed metrics, not on generic football opinions or unconstrained text generation.

2. AI is used as an augmentation layer

LLMs and prompt-driven workflows will be used to translate analytical outputs into structured, useful deliverables. Their role is interpretive and communicative, not foundational.

3. Outputs must remain structured and explainable

Each stage should produce outputs in a form that can be reused by later stages. Wherever possible, intermediate outputs should be saved in structured formats such as tables, dictionaries, or JSON-like records rather than only long-form prose.

4. The system should simulate real analyst workflow

The project is meant to represent how a modern analyst might work in practice: compute metrics, interpret patterns, compare teams, generate recommendations, and produce decision-ready outputs through a repeatable pipeline.

Layer One - Data Ingestion and Preparation

Purpose

This layer is responsible for acquiring, organizing, validating, and standardizing the raw datasets that feed the rest of the project.

Primary Inputs

- SkillCorner Open Tracking Data
- Metrica Sports sample tracking and event data
- Optional event enrichment data such as StatsBomb Open Data

Key Functions

- Load raw tracking and event files
- Standardize file formats and naming conventions
- Normalize coordinate systems and orientation
- Align team and match identifiers
- Validate data completeness and integrity
- Prepare phase-ready datasets for downstream segmentation

Expected Outputs

- cleaned match-level tracking datasets
- standardized coordinate data
- match metadata tables
- optional linked event context tables

Notes

This layer is not where LLMs should play a major role. Cleaning and transformation of numeric tracking data should remain conventional and programmatic. At most, AI can assist with documentation, schema interpretation, or pipeline notes, but not with core numeric cleaning logic.

Layer Two - Phase Segmentation and Tactical Feature Engineering

Purpose

This layer transforms raw tracking data into football-relevant analytical variables. It is where the system begins converting movement data into interpretable tactical structure.

Core Segmentation Contexts

- in possession
- out of possession
- defensive transition
- attacking transition

Possible Feature Families

- defensive line height
- team width and team depth
- compactness measures
- inter-player spacing
- convex hull area
- centroid location and movement
- spacing between lines
- transition recovery speed
- structural stability after possession loss
- pressing reaction proxies
- ball distance relationships
- corridor occupation measures
- role consistency indicators
- synchronization or coordinated movement measure

Expected Outputs

- frame-level metric tables
- possession/transition segment summaries
- match-level tactical profile dataset
- team-level tactical fingerprint dataset

Key Functions

- segment matches by phase of play
- calculate frame-level spatial metrics
- aggregate metrics into sequence-level summaries
- aggregate sequence results into match-level and team-level profiles
- create comparable tactical fingerprints across matches

Notes

This is one of the true analytical cores of the project. The quality of everything downstream depends on whether these features are meaningful, stable, and interpretable.

Layer 3 - Tactical Modeling and Identity Analysis

Purpose

This layer takes engineered features and attempts to identify patterns, tactical identities, consistencies, and stylistic differences across teams or matches.

Key Analytical Goals

- measure a team's tactical identity
- evaluate whether that identity is stable across matches
- compare one team to another
- identify recurring structural strengths and weaknesses
- quantify stylistic similarities and differences
- support later recommendation generation

Possible Analytical Methods

- descriptive profiling
- clustering / unsupervised learning
- similarity scoring
- stability analysis across matches
- variance / entropy analysis
- simple rule-based thresholds for tactical interpretation
- selective dimensionality reduction if needed

Expected Outputs

- cluster labels or style groups
- tactical similarity matrices
- identity stability measures
- match-to-match variation summaries
- flagged strengths / weaknesses based on metric thresholds or relative standing

Notes

This layer should remain restrained. The goal is not to force advanced modeling everywhere. If clustering or dimensionality reduction genuinely adds value, use it. If not, clear descriptive identity profiling may be stronger and easier to defend.

Layer 4 -AI-Assisted Interpretation and Decision Support

This is the major extension that differentiates the project. The AI layer sits on top of the computed analytics. It receives structured metric outputs and uses them to generate consistent, constrained, reusable interpretations and recommendations. This layer should be divided into a small number of clearly defined modules.

Module A - Tactical Interpretation Engine

Purpose

Translate quantitative tactical metrics into structured qualitative interpretation.

Inputs

- team-level tactical fingerprint
- match-level tactical summaries
- selected comparisons to league/sample baselines
- metric definitions and interpretation rules

Key Tasks

- summarize tactical identity
- identify key behavioral tendencies
- highlight strengths and weaknesses
- describe how the team tends to defend, transition, or organize structurally

Example Output Categories

- team identity summary
- strongest tactical traits
- weakest tactical traits
- transition behavior summary
- compactness / structure summary
- confidence notes or caution statements

Expected Structured Output

```
{
    "team_identity_summary": "...",
    "strengths": ["...", "..."],
    "weaknesses": ["...", "..."],
    "notable_patterns": ["...", "..."],
    "confidence_notes": ["..."]
}
```

Role in the system

This module turns metrics into analyst-style interpretation. It is useful because many final users need interpretation, not raw tables.

Module B - Recommendation Engine

Purpose

Generate data-grounded decision-support suggestions from identified strengths, weaknesses, and patterns.

Inputs

- outputs from Tactical Interpretation Engine
- selected raw or aggregated metrics
- optional predefined recommendation rules or prompt constraints

Key Tasks

- suggest tactical adjustments
- identify training priorities
- suggest areas for structural improvement
- translate weaknesses into action-oriented recommendations

Example Output Categories

- tactical adjustments
- training focus recommendations
- structural risk areas
- opponent-prep implications
- analyst follow-up suggestions

Expected Structured Output

```
{
    "risk_areas": ["...", "..."],
    "tactical_adjustments": ["...", "..."],
    "follow_up_questions": ["...", "..."],
    "training_focus_areas": ["...", "..."],
}
```

Important constraint

This module must remain grounded in evidence already surfaced by the analytics. It should not invent tactics or make claims unsupported by the metrics.

Module C - Opponent / Matchup Scouting Engine

Purpose

Compare two teams and generate a structured matchup-oriented interpretation.

Inputs

- Team A tactical fingerprint
- Team B tactical fingerprint
- comparative strengths/weaknesses
- optional similarity and contrast metrics

Key Tasks

- identify stylistic contrasts
- identify where one team may exploit the other
- highlight key structural matchup points
- suggest opponent-prep focus areas

Example Output Categories

- matchup overview
- exploitable tendencies
- defensive warnings
- likely transition battlegrounds
- preparation priorities

Expected Structured Output

```
{
    "team_a_risks": ["...", "..."],
    "matchup_summary": "...",
    "team_a_opportunities": ["...", "..."],
    "prep_priorities": ["...", "..."]
}
```

Role in the system

This module makes the project feel more like a real football operations workflow rather than a pure descriptive research project.

Module D - Automated Reporting Engine

Purpose

Assemble outputs from prior modules into clear, reusable deliverables.

Inputs

- tactical metrics and visuals
- interpretation outputs
- recommendation outputs
- matchup outputs if applicable

Key Tasks

- build tactical summary reports
- build opponent-prep reports
- generate executive-style summaries
- generate technical explanation notes

Example Output Categories

- markdown reports
- notebook outputs
- PDF-ready text blocks
- dashboard-ready text
- slide-ready summaries

Role in the system

This module is the communication layer. It demonstrates business and workflow value because it shows how analysis becomes usable output.

Layer 5 - Final Output and Delivery Layer

Purpose

Present the results of the system in forms that are useful for different audiences.

Possible Final Deliverables

- technical tactical report
- executive summary
- AI workflow appendix
- opponent-prep memo
- tactical identity summary sheet
- lightweight dashboard or visual interface
- portfolio writeup
- GitHub documentation

Audience Types

- technical evaluator
- professor / class audience
- hiring manager
- football operations audience
- general portfolio reviewer

Workflow Sequence

Below is the intended workflow from raw data to final output.

1. Ingest and standardize data

Tracking and event data are loaded, normalized, and validated.

2. Segment phases of play

Possession state and transition contexts are identified so metrics can be computed meaningfully.

3. Compute tactical metrics

Spatial-temporal metrics are calculated and aggregated into sequence-, match-, and team-level forms.

4. Build tactical fingerprints

Computed metrics are organized into profiles that describe teams' structural and behavioral tendencies.

5. Analyze tactical identity

Patterns, consistencies, similarities, and weaknesses are identified using descriptive and/or unsupervised methods.

6. Pass structured outputs into AI layer

Relevant metrics and summaries are formatted into structured inputs for downstream prompt-based modules.

7. Generate tactical interpretations

The Tactical Interpretation Engine produces identity summaries, strengths, weaknesses, and key patterns.

8. Generate recommendations

The Recommendation Engine translates the interpretation into decision-support suggestions.

9. Generate opponent comparison if applicable

The Matchup Scouting Engine compares two teams and produces opponent-prep style outputs.

10. Assemble reports

The Automated Reporting Engine combines all prior outputs into communication-ready forms.

Input / Output Logic Between Modules

This part matters because it makes the system feel engineered rather than vague.

Analytics Layer → Interpretation Layer

The analytics layer should output structured, machine-readable summaries rather than loose notes. These may include:

- metric dictionaries
- ranking summaries
- percentile-like comparisons
- flagged threshold conditions
- descriptive bullet summaries generated programmatically

Interpretation Layer → Recommendation Layer

The recommendation module should not re-read the whole dataset. It should consume:

- structured strengths
- structured weaknesses
- selected key metrics
- confidence notes

Recommendation Layer → Reporting Layer

The reporting layer should receive:

- summary statements
- recommendation blocks
- key figures
- selected charts or visuals

This chaining is important because it demonstrates workflow design and modularity.

Human in the Loop Design

A core rule of the system is that human oversight remains central. The analyst remains responsible for:

- metric design
- data validation
- interpretation review
- prompt design
- checking whether recommendations are reasonable
- deciding which outputs are ultimately used

The AI layer is not an autonomous tactical expert. It is an assistant layer that helps convert validated analytical outputs into faster, more consistent, and more structured decision-support artifacts.

This distinction is important both for project credibility and for honest documentation of AI use.

Prompting and Structured Output Philosophy

To keep the AI layer reliable, prompts should be designed around:

- explicit role definitions
- strict grounding in provided metrics
- output schema constraints
- instructions to avoid unsupported claims
- instructions to express uncertainty where needed
- instructions to separate observation from recommendation

Where possible, prompts should ask for structured outputs first and prose second.

Good:

- list strengths
- list weaknesses
- tie each to specific metric evidence
- generate recommendations from those weaknesses

Bad:

- “Tell me what this team is like”

The project should emphasize that prompt design is part of the system architecture, not an afterthought.

System Boundaries

To prevent scope drift, this system should explicitly avoid a few things unless time remains later.

Not core to version 1

- full chatbot interface
- autonomous multi-agent orchestration
- overbuilt dashboard platform
- advanced deployment infrastructure
- LLM-based numeric cleaning of raw tracking data
- unsupported predictive claims beyond the data

What matters more

- a clean analytics core
- a small number of strong AI modules
- clear inputs and outputs
- believable decision-support outputs
- strong documentation of workflow design

What This Architecture Demonstrates Professionally

This system is valuable because it demonstrates both traditional analytical ability and modern AI workflow thinking. It shows:

- data engineering discipline
- sports analytics reasoning
- tactical metric design
- structured workflow design
- LLM integration in a real use case
- human-in-the-loop AI usage
- communication and reporting system design

What Short Version of the Architecture

This project uses football tracking data to compute tactical identity metrics and structural team profiles, then layers an AI-assisted workflow on top of those outputs to generate tactical interpretations, decision-support recommendations, opponent scouting summaries, and automated reports. The system is designed as a modular, human-in-the-loop pipeline that preserves analytical rigor while demonstrating how AI can be integrated into real modern analytics workflows.